

Mashbot — Defect Reporting and Final Quality Judgment

SE3002 Assignment 01, Part 4

Jira project: **KAN** ("My Software Team"), site `chabdullah3506.atlassian.net`. All 5 issues below were investigated and confirmed reproducible before logging, per the assignment's requirement that a test status and a defect are not the same thing.

Defect Log Summary

Jira	ID	Title	Severity	Priority	Test case	Status
KAN-4	BUG-01	Deactivated account bypasses login block	Critical	Highest	TC-02	To Do (unfixed)
KAN-5	BUG-02	Scheduled go-live defaults to current time instead of 00:00	Medium	Medium	TC-04	To Do (unfixed)
KAN-6	BUG-03	Campaign edit surfaces raw DB error instead of a validation message	Low-Medium	Low	TC-07	To Do (unfixed)
KAN-7	BUG-04	Campaign/admin/accounts action buttons crash the page (RSC boundary violation)	Critical	Highest	TC-15	Done (fixed, PR #3)
KAN-8	BUG-05	RLS policy blocks Contributors from submitting content for approval	Critical	Highest	TC-08	Done (fixed, PR #4)

Each Jira issue contains: environment/build (mashbot2.0, baseline-v2 / commit 74bd4f8a), severity, priority, reproducibility, preconditions, numbered reproduction steps, expected vs. actual result, root cause, related test case ID, and evidence — per the assignment's required fields. Full text is in Jira; issue links: `chabdullah3506.atlassian.net/browse/KAN-4` through `KAN-8`.

Investigation note: BUG-01/02/03 were reproduced multiple times with consistent results and traced to a specific line of source before logging. BUG-04 and BUG-05 were discovered live during Part 3B execution, blocked further testing entirely, and were therefore fixed immediately — as separate commits on top of the untouched baseline-v2 tag, not by altering the evaluated baseline itself.

Final Quality Judgment

The evaluated scope is the 7 FR + 3 NFR baseline frozen at git tag `baseline-v2`, assessed through a real SonarQube Cloud scan and 15 manually executed test cases run against the live, Supabase-backed application — not a mock or offline build.

The evidence supports a mixed but defensible verdict. SonarQube reported 46 issues (1 Security, 45 Maintainability, 0 Reliability) with no automated test coverage configured. The five interpreted findings showed that routine code-smell warnings sometimes point to real defects rather than pure style debt — most notably, an "unused variable" flag traced directly to a genuine login-security failure. Manual testing confirmed this pattern independently: 12 of 15 cases passed, but 5 reproducible defects were found (BUG-01 through BUG-05), spanning access control (a deactivated account can still log in), a business-rule violation (scheduling defaults to the current time, not midnight, contradicting SRS 0530), inconsistent validation between campaign creation and editing, and two defects severe enough — a Next.js server/client boundary crash and a database row-level-security error — that they blocked all further test execution until fixed. All five are logged in

Jira with root causes and reproduction steps.

What remains unsupported is narrower but real. The backup and network-encryption NFRs rely on Supabase/hosting-platform guarantees that cannot be independently verified from application code; the memory-constraint NFR was spot-checked, not exhaustively profiled. Three confirmed defects (BUG-01/02/03) remain unfixed by deliberate choice, preserving the frozen baseline for evaluation — meaning the graded baseline is knowingly non-compliant with SRS 0390, 0530, and its own campaign-date-validation rule.

AI-introduced assumptions materially shaped what could be evaluated at all. Simulated external-service integrations (no real OAuth) meant the external-publish requirement could never be executed as a genuine end-to-end test rather than a documented limitation, and disclosed scoping choices — global roles instead of per-campaign permissions, a fixed session timeout — narrow how much of the original SRS this MVP can fairly be judged against.

On balance, the 10-requirement scope is acceptable as a development-stage baseline, not a release candidate: the core Contributor -> Approver -> Publisher workflow demonstrably works end-to-end once the two blocking defects were fixed, but the three remaining confirmed defects — particularly BUG-01's access-control failure — must be resolved before this could reasonably be called production-acceptable.

(364 words)